

► 注意事项:

- 这~~不是~~是一个逐行代码讲解的视频
- 只讲关键部分的代码实现, 可以一定程度上减轻复现或者读懂代码的难度
- 深入理解DDPM
- DDPM不算是很基础的深度学习内容, 需要有基本的深度学习的知识和Pytorch的使用经验

- ◀
- 比较简单的玩具代码: <https://github.com/abarankab/DDPM> 主要讲解
 - 比较实用的代码: <https://github.com/Janspiry/Image-Super-Resolution-via-Iterative-Refinement>
- ▶

- 为了复现DDPM, 我们需要考虑哪些细节问题?
- (1) alpha、beta这些预先定义好的参数怎么获得?
- (2) 训练过程? ✓
- (3) 推理过程? ✓
- (4) UNet的结构?

α_t β_t

$$\underline{z} = \underline{\text{UNet}}(\underline{x_t}, t)$$

↑
也可以换成别的网络
但Unet开销小

$$\beta_t \ 10^{-2} \rightarrow 2 \times 10^{-4}$$

$$T=2000$$

- (1) alpha、beta这些预先定义好的参数怎么获得?
- <https://github.com/abarankab/DDPM/blob/main/ddpm/diffusion.py>

$\bar{\alpha}_t$ alphas = 1.0 - betas
alphas_cumprod = np.cumprod(alphas) 连乘
to_torch = partial(torch.tensor, dtype=torch.float32)

把 np.array 转化成 tensor

```
self.register_buffer("betas", to_torch(betas))
self.register_buffer("alphas", to_torch(alphas))
self.register_buffer("alphas_cumprod", to_torch(alphas_cumprod))

self.register_buffer("sqrt_alphas_cumprod", to_torch(np.sqrt(alphas_cumprod)))
self.register_buffer("sqrt_one_minus_alphas_cumprod", to_torch(np.sqrt(1 - alphas_cumprod)))
self.register_buffer("reciprocal_sqrt_alphas", to_torch(np.sqrt(1 / alphas)))

self.register_buffer("remove_noise_coeff", to_torch(betas / np.sqrt(1 - alphas_cumprod)))
self.register_buffer("sigma", to_torch(np.sqrt(betas)))
```

变成 tensor
保存起来

```
def generate_linear_schedule(T, low, high):
    return np.linspace(low, high, T)
```

$\alpha_t \leftarrow \textcircled{t}$
索引

- (2) 训练过程?
- <https://github.com/abarankab/DDPM/blob/main/ddpm/diffusion.py>

```
def get_losses(self, x, t, y):
    noise = torch.randn_like(x)

    perturbed_x = self.perturb_x(x, t, noise)
    estimated_noise = self.model(perturbed_x, t, y)

    if self.loss_type == "l1":
        loss = F.l1_loss(estimated_noise, noise)
    elif self.loss_type == "l2":
        loss = F.mse_loss(estimated_noise, noise)

    return loss

def forward(self, x, y=None):
    b, c, h, w = x.shape
    device = x.device

    if h != self.img_size[0]:
        raise ValueError("image height does not match diffusion parameters")
    if w != self.img_size[0]:
        raise ValueError("image width does not match diffusion parameters")

    t = torch.randint(0, self.num_timesteps, (b,), device=device)

    return self.get_losses(x, t, y)
```

退化 X_t
预测 \hat{z}

选择损失函数

训练得到一个 model


原图 x_0

检查图像合规性

任取一个 t

时间最大步

和 batch size 的维度一样



- (3) 推理过程?
- <https://github.com/abarankab/DDPM/blob/main/ddpm/diffusion.py>

采样
方式

```

@torch.no_grad()
def remove_noise(self, x, t, y, use_ema=True):
    if use_ema:
        return (
            (x - extract(self.remove_noise_coeff, t, x.shape) * self.ema_model(x, t, y)) *
            extract(self.reciprocal_sqrt_alphas, t, x.shape)
        )
    else:
        return (
            (x - extract(self.remove_noise_coeff, t, x.shape) * self.model(x, t, y)) *
            extract(self.reciprocal_sqrt_alphas, t, x.shape)
        )

```

去噪 Σ

```

@torch.no_grad()
def sample(self, batch_size, device, y=None, use_ema=True):
    if y is not None and batch_size != len(y):
        raise ValueError("sample batch size different from length of given y")

```

正态分布随机采样 x_T

```

x = torch.randn(batch_size, self.img_channels, *self.img_size, device=device)

```

 x_t

```

for t in range(self.num_timesteps - 1, -1, -1):

```

一共做 T 步上推过程

```

    t_batch = torch.tensor([t], device=device).repeat(batch_size)

```

转换成 tensor

```

    x = self.remove_noise(x, t_batch, y, use_ema)

```

得到 x_{t-1} 的均值部分 x_{t-1}

t > 0

要加噪声

```

    if t > 0:
        x += extract(self.sigma, t_batch, x.shape) * torch.randn_like(x)

```

```

    return x.cpu().detach()

```


- (4) UNet的结构?
- <https://github.com/abarankab/DDPM/blob/main/ddpm/unet.py>

```
x = self.init_conv(x)
```

```
skips = [x]
```

```
for layer in self.downs:
```

```
    x = layer(x, time_emb, y)
```

```
    skips.append(x)
```

```
for layer in self.mid:
```

```
    x = layer(x, time_emb, y)
```

```
for layer in self.ups:
```

```
    if isinstance(layer, ResidualBlock):
```

```
        x = torch.cat([x, skips.pop()], dim=1)
```

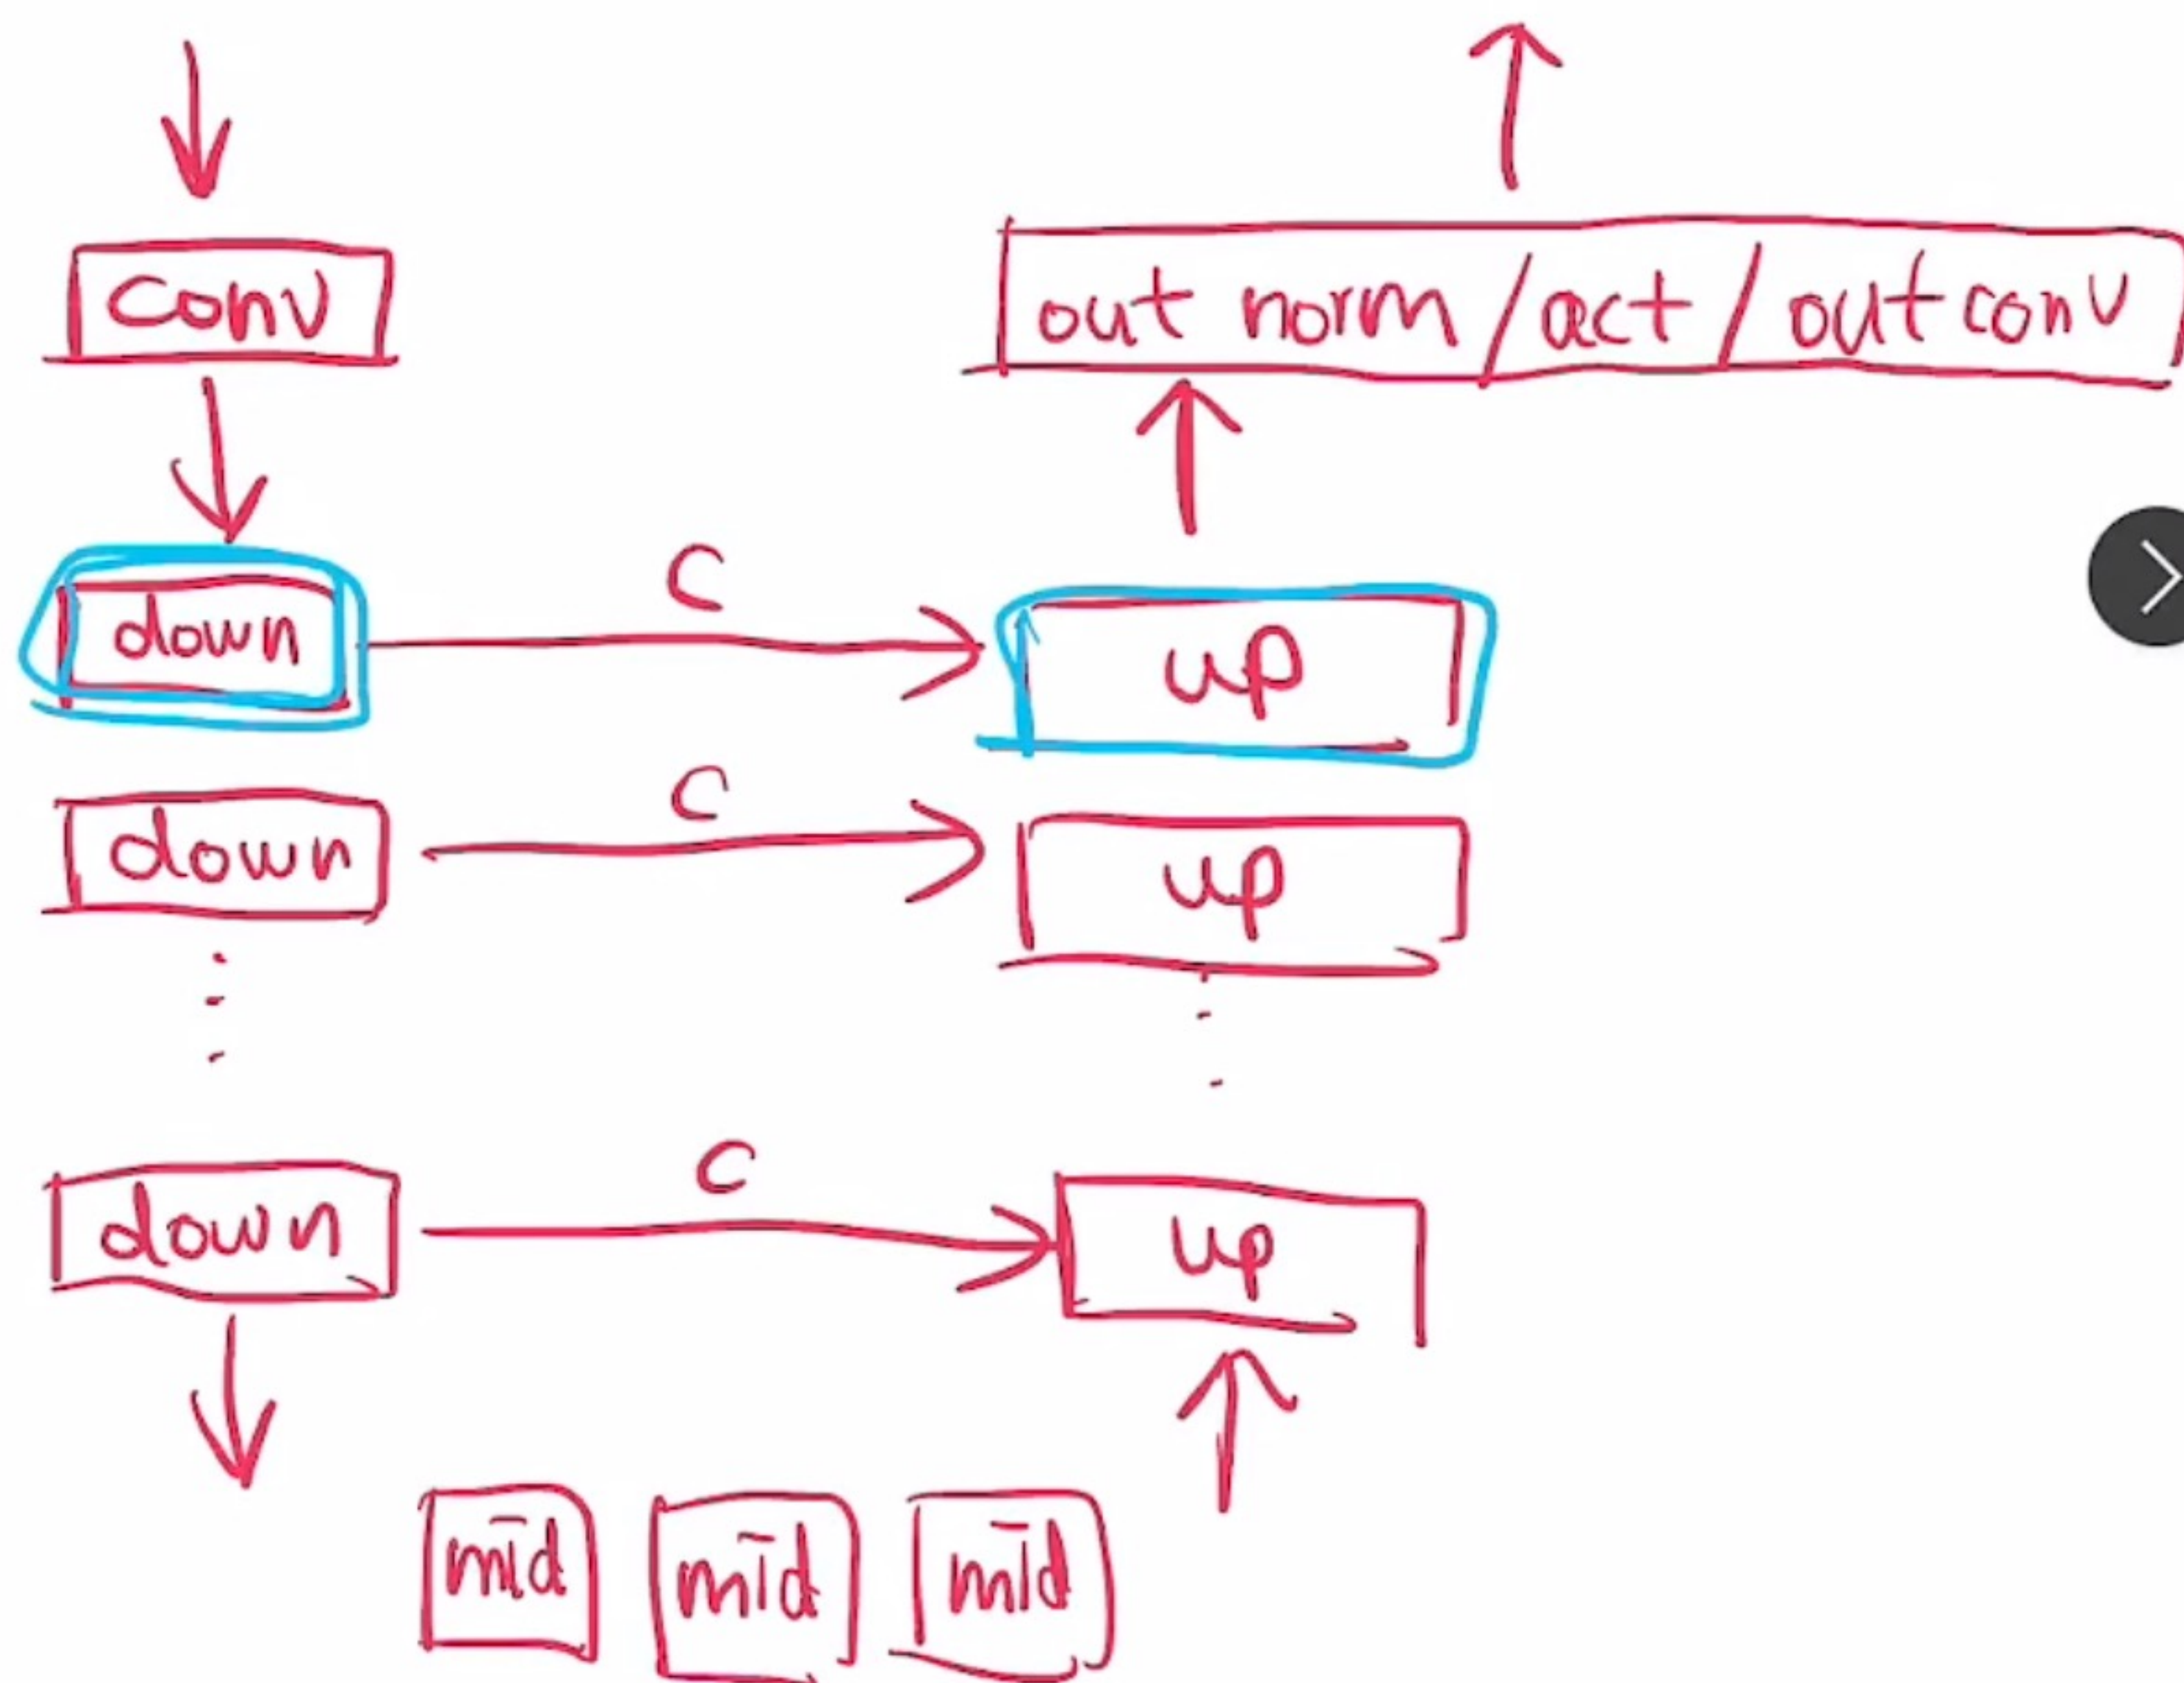
```
        x = layer(x, time_emb, y)
```

```
x = self.activation(self.out_norm(x))
```

```
x = self.out_conv(x)
```

concat

这是关键



➤ (4) UNet的结构?

➤ <https://github.com/abarankab/DDPM/blob/main/ddpm/unet.py>

```
x = self.init_conv(x)
```

```
skips = [x]
```

```
for layer in self.downs:
```

```
    x = layer(x, time_emb, y)
```

```
    skips.append(x)
```

```
for layer in self.mid:
```

```
    x = layer(x, time_emb, y)
```

```
for layer in self.ups:
```

```
    if isinstance(layer, ResidualBlock):
```

```
        x = torch.cat([x, skips.pop()], dim=1)
```

```
    x = layer(x, time_emb, y)
```

```
x = self.activation(self.out_norm(x))
```

```
x = self.out_conv(x)
```

向前去找

```
if self.time_mlp is not None:
```

```
    if time is None:
```

```
        raise ValueError("time conditioning was specified but tim is not passed")
```

```
    time_emb = self.time_mlp(time)
```

```
else:
```

```
    time_emb = None
```

mlp 输出表示时间的向量

t

20 30

0 ~ 2000

➤ (4) UNet的结构?

➤ <https://github.com/abarankab/DDPM/blob/main/ddpm/unet.py>

```
x = self.init_conv(x)
```

```
skips = [x]
```

```
for layer in self.downs:
    x = layer(x, time_emb, y)
    skips.append(x)
```

```
for layer in self.mid:
    x = layer(x, time_emb, y)
```

是怎么送进模型的呢?

```
for layer in self.ups:
    if isinstance(layer, ResidualBlock):
        x = torch.cat([x, skips.pop()], dim=1)
    x = layer(x, time_emb, y)
```


```
x = self.activation(self.out_norm(x))
x = self.out_conv(x)
```

```
if self.time_mlp is not None:
    if time is None:
        raise ValueError("time conditioning was specified but tim is not passed")

    time_emb = self.time_mlp(time)
else:
    time_emb = None
```

```
self.time_mlp = nn.Sequential(
    PositionalEmbedding(base_channels, time_emb_scale),
    nn.Linear(base_channels, time_emb_dim),
    nn.SiLU(),
    nn.Linear(time_emb_dim, time_emb_dim),
) if time_emb_dim is not None else None
```

非线性过程



- (4) UNet的结构?
- <https://github.com/abarankab/DDPM/blob/main/ddpm/unet.py>

```
x = self.init_conv(x)
```

```
skips = [x]
```

```
for layer in self.downs:
```

```
    x = layer(x, time_emb, y)
```

```
    skips.append(x)
```

```
for layer in self.mid:
```

```
    x = layer(x, time_emb, y)
```

```
for layer in self.ups:
```

```
    if isinstance(layer, ResidualBlock):
```


```
        x = torch.cat([x, skips.pop()], dim=1)
```

```
        x = layer(x, time_emb, y)
```

```
x = self.activation(self.out_norm(x))
```

```
x = self.out_conv(x)
```

RB
残差模块



```
def forward(self, x, time_emb=None, y=None):
```

```
    out = self.activation(self.norm_1(x))
```

```
    out = self.conv_1(out)
```

```
    if self.time_bias is not None:
```

```
        if time_emb is None:
```

```
            raise ValueError("time conditioning was specified but time_emb is not passed")
```

```
            out += self.time_bias(self.activation(time_emb))[:, :, None, None]
```

加入
time_emb

```
    if self.class_bias is not None:
```

```
        if y is None:
```

```
            raise ValueError("class conditioning was specified but y is not passed")
```

```
            out += self.class_bias(y)[:, :, None, None]
```

```
    out = self.activation(self.norm_2(out))
```


```
    out = self.conv_2(out) + self.residual_connection(x)
```

```
    out = self.attention(out)
```

自注意力

```
    return out
```


► CIFAR10训练150K次迭代的随机生成结果



建议：

- 想要掌握DDPM，看代码，跑实验，出结果
- 想要看更好的生成结果，参考<https://github.com/Janspiry/Image-Super-Resolution-via-Iterative-Refinement>
- 看不懂代码？先跑实验再说

更新计划（大概吧）：

- 第一章：DDPM介绍 ←
 - 第二章：DDPM的应用
 - 第三章：score-based model介绍 ←
 - 第四章：score-based model的应用
- SR3 debiut DDIM Diffusion CLIP ILVR
SPedit MRI ...

扩散模型

score-based ✓

偏理论多

DDPM ✓

偏应用多